

atividade4

April 22, 2025

0.1 Atividade Curta 4 - Treinamento de rede neural usando batch

Aluno: 119891 - João Vitor Silva de Oliveira Test with Iris dataset and Test and train split function

```
[1]: from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split

iris = load_iris()
X, y = iris.data, iris.target

X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=1)

print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)

print(X_train[0])
print(y_test[0])
```

```
(120, 4)
(30, 4)
(120,)
(30,)
[6.1 3.  4.6 1.4]
0
```

0.1.1 Architecture definition

```
[2]: #define ANN architecture as a Torch NN Module

import torch.nn as nn

class ANN(nn.Module):
    def __init__(self, num_classes=3):
        super().__init__()
        self.lin1 = nn.Linear(4, 10)
        self.act1 = nn.ReLU()
```

```

self.lin2 = nn.Linear(10, 20)
self.act2 = nn.ReLU()
self.lin3 = nn.Linear(20, num_classes)

def forward(self, x, debug=False):
    if debug : print(f'Shape de entrada: {x.shape}')
    x = self.lin1(x)
    if debug : print(f'Shape de entrada: {x.shape}')
    x = self.act1(x)
    if debug : print(f'Shape de entrada: {x.shape}')
    x = self.lin2(x)
    if debug : print(f'Shape de entrada: {x.shape}')
    x = self.act2(x)
    if debug : print(f'Shape de entrada: {x.shape}')
    y = self.lin3(x)
    if debug : print(f'Shape de entrada: {y.shape}')
    return y

```

0.1.2 Model analysis

```

[3]: import torch

if torch.cuda.is_available():
    my_device = torch.device("cuda:0")
else:
    my_device = torch.device("cpu")

print(f"Running on {my_device.type}.")

net = ANN( num_classes=3 )
#net = ANN()

net = net.to(my_device)

a = torch.rand( (1, 4) )
a = a.to(my_device)
b = net( a , debug=True)

del a, b, net

```

Running on cuda.

```

Shape de entrada: torch.Size([1, 4])
Shape de entrada: torch.Size([1, 10])
Shape de entrada: torch.Size([1, 10])
Shape de entrada: torch.Size([1, 20])
Shape de entrada: torch.Size([1, 20])
Shape de entrada: torch.Size([1, 3])

```

```
[4]: #!/pip install torchsummary
from torchsummary import summary

net = ANN( num_classes=3 )

net = net.to(my_device)

summary(net, input_size=(1, 4), batch_size=1)

del net
```

```
-----
          Layer (type)                Output Shape          Param #
=====
          Linear-1                    [1, 1, 10]             50
          ReLU-2                      [1, 1, 10]              0
          Linear-3                    [1, 1, 20]            220
          ReLU-4                      [1, 1, 20]              0
          Linear-5                    [1, 1, 3]              63
=====

Total params: 333
Trainable params: 333
Non-trainable params: 0
-----

Input size (MB): 0.00
Forward/backward pass size (MB): 0.00
Params size (MB): 0.00
Estimated Total Size (MB): 0.00
-----
```

0.1.3 Training functions

```
[5]: #!/pip install tensorboard

from datetime import datetime

from torch.utils.tensorboard import SummaryWriter

import matplotlib.pyplot as plt
import copy
from tqdm import tqdm

def train ( dataset, prefix=None, upper_bound=101.0, save=False, epochs=100,
            lr=1e-1, device='cpu', debug=False, layers2tensorboard=True,
            ↪lambda_reg=0 ) :
```

```

num_classes = 3

tensorboard_path = './tensorboard/'

net = ANN( num_classes )
net.to(device)

optimizer = torch.optim.SGD(net.parameters(), lr=lr,
↪weight_decay=lambda_reg)
criterion = nn.CrossEntropyLoss()

now = datetime.now()
suffix = now.strftime("%Y%m%d_%H%M%S")
prefix = suffix if prefix is None else prefix + '-' + suffix

writer = SummaryWriter( log_dir=tensorboard_path+prefix )

accuracies = []
max_accuracy = -1.0

data_loader = train_test_split( dataset.data,
                                dataset.target,
                                test_size=0.2,
                                random_state=1 )

train_x, test_x, train_label, test_label = data_loader

train_x = torch.from_numpy(train_x).float()
train_x = train_x.to(device)
train_label = torch.from_numpy(train_label).float()
train_label = train_label.to(device)

test_x = torch.from_numpy(test_x).float()
test_x = test_x.to(device)
test_label = torch.from_numpy(test_label).float()
test_label = test_label.to(device)

writer.add_graph(net, train_x)

for epoch in tqdm( range(epochs) , desc='Training epochs...' ) :

    # Set Pytorch variables
    net.train()
    optimizer.zero_grad()

    # Forward step
    predict_y = net( train_x )

```

```

# Loss
error = criterion( predict_y , train_label.long() )

# Back propagation
error.backward()
optimizer.step()

# Accuracies:
predict_ys = torch.max( predict_y, axis=1 )[1]
correct     = torch.sum( predict_ys == train_label )
accuracy_train = correct/train_x.size(0)

accuracy_test = validate(net, test_x, test_label, device=device)
accuracies.append(accuracy_test)

# Tensor board writing
writer.add_scalar( 'Loss/train', error.item(), epoch )
writer.add_scalar( 'Accuracy/train', accuracy_train, epoch )
writer.add_scalar( 'Accuracy/test', accuracy_test, epoch )

if layers2tensorboard :
    plot_layers( net, writer, epoch )

# Test model
if accuracy_test > max_accuracy:
    best_model = copy.deepcopy(net)
    max_accuracy = accuracy_test
    print(f'Saving the best model at epoch {epoch+1:3d} ' +
          f'with Accuracy: {accuracy_test:8.4f}%')

if debug : print( f'Epoch: {epoch+1:3d} |'
                  + f'Accuracy Test: {accuracy_test:3.4f}%' )

if accuracy_test > upper_bound :
    break

if save :
    models_path = './models/'
    path = f'{models_path}{prefix}-{max_accuracy:.2f}.pkl'
    torch.save( best_model, path )
    print( f'Model saved in: {path}' )

plt.figure(figsize=(16, 8))
plt.plot(accuracies)

writer.flush()

```

```
writer.close()

return best_model
```

```
[6]: def validate ( model , test_x, test_label , device='cpu') :

    model.eval()

    predict_y = model( test_x ).detach()
    predict_ys = torch.max( predict_y, axis=1 )[1]
    correct = torch.sum(predict_ys == test_label)

    return correct.to('cpu').numpy()*100./test_x.size(0)
```

```
[7]: def plot_layers ( net , writer, epoch ) :
    layers = list(net.modules())

    layer_id = 1
    for layer in layers:
        if isinstance(layer, nn.Linear) :
            writer.add_histogram(f'Bias/linear-{layer_id}', layer.bias, epoch )
            writer.add_histogram(f'Weight/linear-{layer_id}', layer.weight, epoch )
            writer.add_histogram(f'Grad/linear-{layer_id}', layer.weight.grad, epoch )
            layer_id += 1
```

0.1.4 Run the training phase

```
[8]: from sklearn.datasets import load_iris

if torch.cuda.is_available():
    my_device = torch.device("cuda:0")
else:
    my_device = torch.device("cpu")

print(f"Running on {my_device.type}")


dataset = 'Iris'
epochs = 500
lr = 1e-1
lambda_reg = 1e-4
prefix = 'ANN-{}-e-{}-lr-{}'.format(dataset, epochs, lr)

iris = load_iris()
dataset = iris
```

```
net = train( dataset=dataset, epochs=epochs, device=my_device,
             upper_bound=100.0, lr=lr, lambda_reg=lambda_reg,
             layers2tensorboard=True, save=False, prefix=prefix
           )
```

Running on cuda

```
Training epochs...:  2%|          | 8/500 [00:00<00:12, 39.05it/s]
Saving the best model at epoch  1 with Accuracy:  43.3333%
Saving the best model at epoch 11 with Accuracy:  53.3333%
Saving the best model at epoch 12 with Accuracy:  56.6667%
Training epochs...:  5%|          | 26/500 [00:00<00:08, 52.80it/s]
Saving the best model at epoch 23 with Accuracy:  60.0000%
Saving the best model at epoch 27 with Accuracy:  96.6667%
Training epochs...: 53%|          | 263/500 [00:05<00:04, 48.67it/s]
Saving the best model at epoch 255 with Accuracy: 100.0000%
Training epochs...: 100%|         | 500/500 [00:10<00:00, 48.59it/s]
```



```
[9]: print(net)

import numpy as np

test_x = torch.from_numpy(X_test).float()
test_x = test_x.to(my_device)
test_label = torch.from_numpy(y_test).float()
test_label = test_label.to(my_device)
```

```

indice = np.random.randint(0,30)

print(test_x[indice])
output = net(test_x[indice])
print(output)
predictedd_class = torch.max(output, dim=0)[1]

print(f'saída do modelo: {predictedd_class.to('cpu').numpy()} e saída correta_
↳{y_test[indice]} para o índice {indice}')

```

```

ANN(
  (lin1): Linear(in_features=4, out_features=10, bias=True)
  (act1): ReLU()
  (lin2): Linear(in_features=10, out_features=20, bias=True)
  (act2): ReLU()
  (lin3): Linear(in_features=20, out_features=3, bias=True)
)
tensor([6.9000, 3.1000, 5.1000, 2.3000], device='cuda:0')
tensor([-6.2318,  2.8725,  3.7826], device='cuda:0', grad_fn=<ViewBackward0>)
saída do modelo: 2 e saída correta 2 para o índice 6

```

1 Training with Batches

```

[10]: from torch.utils.data import DataLoader, TensorDataset

def train_with_batches ( dataset, prefix=None, upper_bound=101.0, save=False,
↳epochs=100,
    lr=1e-1, device='cpu', debug=False, layers2tensorboard=True,
↳lambda_reg=0, batch_size=16 ) :

    num_classes = 3

    tensorboard_path = './tensorboard/'

    net = ANN( num_classes )
    net.to(device)

    optimizer = torch.optim.SGD(net.parameters(), lr=lr,
↳weight_decay=lambda_reg)
    criterion = nn.CrossEntropyLoss()

    now = datetime.now()
    suffix = now.strftime("%Y%m%d_%H%M%S")
    prefix = suffix if prefix is None else prefix + '-' + suffix

```



```

writer = SummaryWriter( log_dir=tensorboard_path+prefix )

accuracies = []
max_accuracy = -1.0

data_loader = train_test_split( dataset.data,
                                dataset.target,
                                test_size=0.2,
                                random_state=1 )

train_x,test_x,train_label,test_label = data_loader

train_x = torch.from_numpy(train_x).float()
train_x = train_x.to(device)
train_label = torch.from_numpy(train_label).long()
train_label = train_label.to(device)

test_x = torch.from_numpy(test_x).float()
test_x = test_x.to(device)
test_label = torch.from_numpy(test_label).long()
test_label = test_label.to(device)

train_dataset = TensorDataset(train_x, train_label)
train_loader = DataLoader(train_dataset, batch_size=batch_size,
↪shuffle=True)

for epoch in tqdm( range(epochs) , desc='Training epochs...' ) :

    net.train()
    epoch_loss = 0.0
    correct = 0
    total = 0

    for batch_x, batch_y in train_loader:
        batch_x, batch_y = batch_x.to(device), batch_y.to(device)

        optimizer.zero_grad()
        outputs = net(batch_x)
        loss = criterion(outputs, batch_y)
        loss.backward()
        optimizer.step()

        epoch_loss += loss.item() * batch_x.size(0)

        _, predicted = torch.max(outputs, 1)
        correct += (predicted == batch_y).sum().item()
        total += batch_y.size(0)

```

```

accuracy_train = correct / total

accuracy_test = validate(net, test_x, test_label, device=device)
accuracies.append(accuracy_test)

# Tensor board writing
writer.add_scalar( 'Loss/train', epoch_loss/total, epoch )
writer.add_scalar( 'Accuracy/train', accuracy_train, epoch )
writer.add_scalar( 'Accuracy/test', accuracy_test, epoch )

if layers2tensorboard :
    plot_layers( net, writer, epoch )

# Test model
if accuracy_test > max_accuracy:
    best_model = copy.deepcopy(net)
    max_accuracy = accuracy_test
    print(f'Saving the best model at epoch {epoch+1:3d} ' +
          f'with Accuracy: {accuracy_test:8.4f}%')

if debug : print( f'Epoch: {epoch+1:3d} |'
                  + f'Accuracy Test: {accuracy_test:3.4f}%' )

if accuracy_test > upper_bound :
    break

if save :
    models_path = './models/'
    path = f'{models_path}{prefix}-{max_accuracy:.2f}.pkl'
    torch.save( best_model, path )
    print( f'Model saved in: {path}' )

plt.figure(figsize=(16, 8))
plt.plot(accuracies)

writer.flush()
writer.close()

return best_model

```

```

[11]: if torch.cuda.is_available():
        my_device = torch.device("cuda:0")
    else:
        my_device = torch.device("cpu")

```

```

print(f"Running on {my_device.type}")

dataset = 'Iris'
epochs = 500
lr = 1e-1
lambda_reg = 1e-4
prefix = 'BS-ANN-{}-e-{}-lr-{}'.format(dataset, epochs, lr)

iris = load_iris()
dataset = iris

net = train_with_batches( dataset=dataset, epochs=epochs, device=my_device,
    upper_bound=101.0, lr=lr, lambda_reg=lambda_reg,
    layers2tensorboard=True, save=False, prefix=prefix, batch_size=16
)

```

Running on cuda

Training epochs...: 1% | 4/500 [00:00<00:14, 33.08it/s]

Saving the best model at epoch 1 with Accuracy: 56.6667%

Saving the best model at epoch 2 with Accuracy: 60.0000%


Saving the best model at epoch 4 with Accuracy: 73.3333%

Saving the best model at epoch 5 with Accuracy: 96.6667%

Training epochs...: 5% | 24/500 [00:00<00:13, 35.22it/s]

Saving the best model at epoch 19 with Accuracy: 100.0000%

Training epochs...: 100% | 500/500 [00:14<00:00, 35.00it/s]



2 Analysis of Results

From the analysis of the results, it is evident that training using batches achieved significantly faster convergence compared to training with the entire dataset. While the model trained with batches reached 100% test accuracy as early as epoch 19, the model trained on the full dataset required 255 epochs to achieve the same performance. This difference highlights how the stochastic noise introduced by batches can accelerate the optimization process and promote earlier generalization, even with a slight loss in gradient precision per epoch. Therefore, in terms of learning efficiency, batch training proved to be considerably more advantageous in this scenario.